

Unified Audio-Visual Latent Space Diffusion: Engineering Native Synesthesia in Generative AI

The Neurobiology and Computation of Sensory Synesthesia

Generative artificial intelligence has historically treated sensory modalities as isolated, independent streams, leading to a fragmented approach to content creation.[1, 2, 3] Traditional pipelines synthesize video and audio sequentially, typically generating silent video frames first and subsequently appending an acoustic track through specialized video-to-audio models.[1, 2, 4] Alternatively, some frameworks operate in reverse, leveraging text-to-audio models followed by audio-to-video synthesis.[2] This cascading methodology introduces profound challenges, including temporal drift, severe lip-sync mismatches, semantic misalignment, and cumulative degradation across generation stages.[5, 6]

In cognitive science, biological synesthesia represents a neurological phenomenon where stimulation of one sensory pathway triggers involuntary, simultaneous experiences in a second, concurrent pathway.[7, 8, 9] Neurological analyses identify structural connections within subcortical and cortical multisensory areas—such as the colliculi and cuneus—as the physical substrates facilitating this immediate cross-modal activation.[10] Furthermore, the Stochastic Resonance Model of synesthesia suggests that elevated levels of neural noise within primary sensory regions allow pre-existing multisensory pathways to achieve the critical thresholds necessary to trigger unified, multi-sensory experiences.[9]

Biological Synesthesia Mechanism:

Sensory Stimulation → Neural Noise Injection → Multisensory Co-Activation

Artificial Synesthesia (Unified Latent Diffusion):

Multimodal Prompt → Co-Denoising Step → Unified Output

Translating this biological blueprint into generative architectures has motivated the development of unified latent space diffusion models.[1, 11] Rather than treating vision, time, and audition as separate entities, modern systems fuse these dimensions within a shared latent continuum.[11] This design allows the model to capture the deep, underlying physics of the physical world, wherein events do not merely possess visual or auditory properties, but exist as coupled spatiotemporal and acoustic phenomena.[7, 11] By co-denoising noisy representations of multiple modalities in a single, synchronized reverse diffusion process, these models establish a mathematical equivalent of native synesthesia, producing highly synchronized, high-fidelity outputs.[1, 6, 11]

This cross-sensory triggering mechanism has been adapted to other physical domains.[7, 8] The Synesthesia of Vehicles framework predicts tactile excitations of vehicle tires directly from forward-looking road images.[7] By establishing a synesthetic mapping between visual and tactile modalities, the model eliminates spatiotemporal gaps, allowing autonomous systems to anticipate road excitations and optimize dynamic vehicle control.[7] Similarly, the Robot Synesthesia framework connects a robotic painter’s action space directly to user-driven sonic interactions.[8] By encoding simulated paintings and input sounds into the same latent space, the system allows the audio to drive the content of the painting.[8] For speech guidance, the system decouples audio into text transcripts and emotional tone, estimating emotions from the tone via a pre-trained canvas-to-emotion model compared with speech emotion recognition modules to guide the artistic mood of the brush strokes.[8]

Mathematical Framework of Spatiotemporal and Acoustic Latent Spaces

To bypass the extreme computational demands of processing raw, high-resolution pixel grids and high-frequency audio waveforms, unified models operate in compressed, continuous latent spaces.[1, 12, 13] These spaces are established using pre-trained, high-fidelity autoencoders.[1, 13, 14]

Let E
 a

and D
 a

denote the audio autoencoder, and let z
 0
 a

$=E$
 a

$(x$
 a
 $)$ represent the clean audio latent representation derived from a raw audio spectrogram x
 a

[14] In the forward diffusion process, Gaussian noise is systematically added to the clean latent across a linear or cosine noise schedule parameterized by β
 t

[1, 14]:
 $q(z$

t
a

|z
t-1
a

)=N(
1-β
t

z
t-1
a

,β
t

l)
q(z
t
a

|z
0
a

)=N(
α
-
t

z
0
a

,(1-
α
-
t

)l)

where

α

-

t

$= \prod$

$i=1$

t

$(1-\beta$

i

).[14] The training objective of the audio denoiser network ϵ

θ

a

, conditioned on a cross-modal prompt x

p

containing text, visual, or other sensory conditions, minimizes the simplified variational lower bound [14]:

L

LDM

$=E$

z

0

a

$, \epsilon, t$

$[\parallel \epsilon - \epsilon$

θ

a

$(z$

t

a

$, t, x$

p

$) \parallel$

2

2

]

During inference, the model recovers the clean audio latent step-by-step.[14] The predicted clean audio latent at any given timestep t is expressed as [14]:

z

$\sim$

0

a

$=$

α

$-$

t

1

$(z$

t

a

$-$

$1-$

α

$-$

t

ϵ

θ

a

$(z$

t

a

$,t,x$

p

$))$

The final acoustic waveform is reconstructed by passing this latent representation through the decoder,

x

$\wedge$

a

$=D$

a

(

z

~

0

a

), and converting the resulting spectrogram via a vocoder.[14]

For joint architectures that transition from diffusion to flow matching, such as MM-Sonate, the model optimizes a vector field v

t

that directly interpolates between the source distribution $p(z$

0

) and target distribution $q(z$

1

).[15, 16, 17] The flow-matching regression objective is defined as:

L

FM

$=E$

$t, q(z$

1

), $p(z$

0

)

$[// v$

t

$(z$

t

)-(z

1

-z

0

)//
2
]

This mathematical formulation avoids the curved trajectories of standard diffusion schedules, enabling faster inference and tighter synchronization between continuous modalities.[15, 17] To model videos, models expand the standard 2D image latent space structured as (Channels×Height×Width) into a 3D spatiotemporal latent structure representing (Channels×Time×Height×Width).[11] For example, Stable Video Diffusion uses a 3D Variational Autoencoder (3D VAE) that compresses video inputs along both spatial and temporal dimensions to preserve visual fidelity while enabling efficient representation learning.[1] NVIDIA's VideoLDM extends this paradigm to high-resolution video generation by pre-training a latent diffusion model on images and subsequently introducing a temporal dimension to the latent space.[12] During temporal decoder fine-tuning, video sequences are processed with a frozen per-frame encoder while enforcing temporally coherent reconstructions across frames, utilizing an additional video-aware discriminator.[12] To achieve high-resolution generation, sparse keyframes are generated at low frame rates, which are then temporally upsampled twice by another interpolation latent diffusion model before spatial upsamplers are applied.[12] This architecture scales to 4.1 billion parameters, with 2.7 billion parameters explicitly trained on videos.[12]

To ensure compression of the acoustic domain, frameworks like Music2latent employ an end-to-end consistency autoencoder that encodes complex-valued Short-Time Fourier Transform (STFT) spectrograms into highly compressed continuous latent spaces.[1] This architecture utilizes cross-connections at all hierarchical levels to condition the consistency model on intermediate encoder outputs, preserving fine-grained acoustic features.[1] Furthermore, to maintain consistency in talking head synthesis, models integrate a pre-trained Speech Autoencoder (SpeechAE) that generates temporally compressed speech latent codes directly aligned with a spatiotemporal highly compressed video latent space learned via a temporal VAE.[18] These synchronized representations are modeled by an Audio-to-Video Diffusion Transformer (A2V-DiT) backbone, which utilizes asynchronous add-noise and asynchronous motion-guided generation in the latent space.[18]

Model / Framework	Video Latent Encoder (Spatio-Temporal Compression)	Audio Latent Encoder (Acoustic Compression)	Alignment / Fusion Method	Primary Source
----------------------	---	--	------------------------------	-------------------

Klear	3D VAE (3 Hz temporal rate; 16× spatial height and width compression) [5]	Audio-VAE (44.1 kHz input waveform downsampled to 43 Hz continuous embeddings, 1024× downsample ratio) [5]	Single-Tower Unified DiT with Omni-Full Attention [5, 6]	[5, 6]
MM-Sonate	3D Causal VAE (~3 Hz temporal rate) [16]	Mel-VAE (43 Hz temporal rate) [16]	Multimodal DiT (MM-DiT) with Flow-Matching [16, 17]	[15, 16, 17]
Veo 3	3D VAE (Channels×Time×Height×Width) [11]	Spectrogram / Waveform fused directly into 3D latent volume [11]	3D U-Net with Audio-Modulated Channels & Skip Connections [11]	[11]
AV-DiT	Flattened video patch embeddings with temporal adapters [13]	2D Mel-Spectrogram via STFT passed through pre-trained VAE [13]	Shared DiT Backbone with Modality Adapters & Cross-Attention [13]	[13, 19]

Structural Paradigms of Generative Audiovisual Networks

The architectural evolution of joint audio-video generation can be categorized into three major paradigms: dual-tower models with shallow interaction, single-tower unified transformers, and training-free inference-time alignment frameworks.[1, 2, 5, 6, 14]

Dual-Tower Architectures and Parameter-Efficient Adaptations

Early frameworks, such as MM-Diffusion, utilize a sequential multimodal dual U-Net design.[1] This approach maintains separate denoising autoencoders for audio spectrograms (using 1D dilated convolutions) and video frames (using decomposed 2D+1D spatial-temporal convolutions).[1] Modality interaction is achieved via cross-connections and random-shift attention blocks that bridge the subnets at each diffusion timestep, enabling each modality's denoising process to attend dynamically to the other.[1, 4]

To optimize computational efficiency, the Audio-Visual Diffusion Transformer (AV-DiT) leverages a shared Diffusion Transformer (DiT) backbone pre-trained solely on static 2D images.[1, 13, 19] Instead of updating all model parameters, AV-DiT freezes the pre-trained image weights and introduces highly localized, trainable adaptations.[4, 13] The video branch incorporates a trainable temporal attention layer into the frozen pre-trained DiT block for temporal consistency, permuting output features into temporal sequences to learn dependencies among frames.[13] To adapt the image-based DiT block for audio generation, AV-DiT integrates audio-specific LoRA layers and dedicated adapters to mitigate the large domain gap and learn the harmonic structure of acoustic spectrograms.[13] Finally, an extra shared block equipped with lightweight parameters facilitates feature interaction between audio and visual modalities by shifting the attention mechanism from single-domain connections to cross-domain interactions, ensuring strict alignment.[13]

This adapter-centric design allows the model to achieve state-of-the-art performance on joint generation benchmarks with significantly fewer tunable parameters, demonstrating that a single shared image generative backbone with modality-specific adaptations is sufficient for constructing a joint audio-video generator.[4, 13, 19] Similarly, the AVControl framework adapts the joint audio-visual foundation model LTX-2 by training each control modality as a separate LoRA on a parallel canvas, providing the reference signal as additional tokens in the attention layers without requiring structural backbone modifications.[19]

Single-Tower Unified Transformers

While dual-tower models reduce computational cost, they often struggle with tight cross-modal coupling, such as speech-to-lip synchronization.[5, 6] To address this limitation, state-of-the-art models employ single-tower architectures.[5, 6]

The Klear framework utilizes unified DiT blocks built around an Omni-Full Attention mechanism.[5, 6] Rather than routing features through separate pathways, each block jointly processes four distinct input streams: raw video tokens, video-related text captions, raw audio

tokens, and audio-related text captions.[5] By performing self-attention over the combined set of multimodal tokens, the model establishes deep cross-modal coupling directly within the hidden layers of the transformer, achieving a higher scaling ceiling and stronger alignment than previous dual-tower systems.[5, 6] The model comprises 26 billion parameters with a flow-matching feed-forward dimension of 4096, incorporating 32 joint diffusion transformer layers combined with multimodal Rotary Position Embeddings (RoPE).[5] Similarly, MM-Sonate adopts a Multi-Modal Diffusion Transformer (MM-DiT) optimized via flow matching.[16, 17] To enforce strict linguistic and temporal coherence during speech generation, MM-Sonate processes a unified instruction-phoneme input.[15, 16, 17] Phonemes extracted via Grapheme-to-Phoneme (G2P) conversion are embedded alongside high-level audio descriptors and scene-level video instructions.[16, 17] All textual conditioning is embedded via a shared text encoder (Qwen2.5-7B for instructions and a dedicated 512-dimensional encoder for phonemes), allowing the single-tower transformer to simultaneously model phonetic durations, acoustic dynamics, and corresponding facial kinematics.[16, 17]

Contrastive Alignment and Guidance Mechanisms

Where training a unified foundation model from scratch is computationally prohibitive, researchers employ inference-time optimization and pre-trained contrastive representations.[2, 14, 20] The Diffusion Latent Aligner (DLA) utilizes a pre-trained ImageBind model to connect isolated, single-modality generative models.[2, 3] Operating similarly to classifier guidance, DLA projects the noisy intermediate latents of the video and audio generators into the shared semantic space of ImageBind.[2, 3] By computing gradients on the directional cosine similarity between the modalities and backpropagating these signals into the respective denoising loops, DLA enforces bidirectional cross-modal alignment without requiring any model fine-tuning.[2, 3] For video-to-audio (V2A) synthesis, the Multimodal Diffusion Guidance (MDG) mechanism enforces semantic consistency through a training-free tri-modal objective.[14] MDG calculates the multidimensional volume spanned by the text, visual, and candidate audio embeddings in a shared latent space.[14] During the reverse diffusion steps of a pre-trained audio model, MDG uses this volume-based alignment metric to steer the denoising trajectory towards semantic alignment with the visual prompt.[14]

Furthermore, Contrastive Audio-Visual Pretraining (CAVP) is widely used to capture fine-grained correlations for synchronized generation.[20, 21, 22] Pioneered in Diff-Foley, CAVP learns semantically and temporally aligned representations by maximizing the contrastive similarity of paired visual and auditory features.[20, 21, 22] A latent diffusion model is then trained on spectrogram latents, using the CAVP-aligned visual features as cross-attention targets.[20, 21, 22] To further ensure temporal synchronization, Diff-Foley employs a double guidance strategy during sampling, simultaneously applying Classifier-Free Guidance (CFG) for audio realism and standard Classifier Guidance (CG) guided by a synchronization classifier.[21, 22]

This contrastive framework has been extended via Generative Multi-Scale CAVP (GMS-CAVP), which uses hierarchical pyramidal pooling and multi-resolution convolutions to align audio and video features across different spatial and temporal scales.[23] The quantitative impact of these architectural elements and visual feature configurations on temporal synchronization accuracy is detailed below in the experimental evaluations of the Diff-Foley system.

Architecture / Model	Primary Core Backbone	Key Alignment / Coupling Mechanism	Zero-Shot & Reference Capabilities	Relative Training & Compute Cost	Primary Source
MM-Diffusion	Dual U-Net Subnets [1]	Interleaved random-shift attention blocks [1]	None (Unconditional generation within specific training domains) [2, 24]	High (Requires training entire dual network from scratch) [4, 24]	[1, 24]
AV-DiT	Modality-shared DiT (frozen ImageNet pre-trained backbone) [4, 13]	Trainable spatial-temporal adapters, audio LoRA, and shared self-attention [13]	Adapts static image priors for synchronized joint audio-video generation [4, 13]	Low (Only newly inserted adapters/LoRA layers are trained) [4, 13]	[4, 13, 19]
Klear	Single-tower DiT (26B parameters, 32 joint layers) [5]	Omni-Full Attention over four concatenated input token streams [5, 6]	High-fidelity instruction-following in both joint and single-modality tasks [6, 15]	Extremely High (Requires large-scale multi-stage training) [5, 6]	[5, 6, 15]

MM-Sonate	Multi-Modal DiT (MM-DiT) with Flow-Matching [16, 17]	Unified phoneme-instruction sequences and multi-head cross-attention [16, 17]	Zero-shot voice cloning and timbre-preserving joint speech synthesis [15, 16, 17]	Extremely High (Trained on 100M aligned audio-video pairs) [25]	[15, 16, 17]
Diff-Foley	Latent Diffusion Model (Spectrogram LDM) [20, 21]	Contrastive Audio-Visual Pretraining (CAVP) and double guidance (CFG + CG) [20, 21, 22]	Generates highly synchronized audio from silent video inputs [20, 21]	Medium (Two-stage training of CAVP encoder and LDM decoder) [20, 22]	[20, 21, 22]
Diffusion Latent Aligner (DLA)	Shared pre-trained single-modality models [2, 3]	ImageBind semantic space projection with bidirectional guidance [2, 3]	Connects arbitrary, off-the-shelf single-modality models [2, 3]	Minimal (Training-free, optimization is performed during inference) [3]	[2, 3]

Quantitative Performance Metrics of Aligners

The evaluation of visual features within the Diff-Foley framework demonstrates that CAVP-aligned visual representations significantly outperform standard backbones in temporal synchronization and audio-visual relevance.[20]

Visual Feature Backbone	Inception Score (IS) ↑	Fréchet Inception Distance (FID) ↓	Kullback-Leibler (KL) Divergence ↓	Alignment Accuracy (Align Acc) % ↑	Primary Source
ResNet50	29.48	17.34	7.34	54.88%	[20]
CLIP	57.43	13.09	5.88	79.23%	[20]
CAVP (Ours)	52.07	11.61	6.33	92.35%	[20]

To evaluate the impact of temporal splits during training, experimental studies within the same framework demonstrate that Temporal Split & Merge Augmentation provides a crucial stabilizing effect during joint generation.[20]

Temporal Augmentation	Inception Score (IS) ↑	Fréchet Inception Distance (FID) ↓	Kullback-Leibler (KL) Divergence ↓	Alignment Accuracy (Align Acc) % ↑	Primary Source
Without Augmentation (×)	50.62	12.65	6.38	79.79%	[20]
With Augmentation (✓)	52.07	11.61	6.33	92.35%	[20]

Advanced Training Strategies and Multi-Scene Data Curation

Training unified models to generate high-fidelity, synchronized audio and video presents a significant challenge: joint optimization often leads to modality imbalance, where the model prioritizes the optimization of one modality (typically vision due to its higher dimensionality) at the expense of the other, resulting in unimodal performance degradation or collapse.[5, 6]

Progressive Multi-Task Curriculum

To stabilize joint optimization and prevent performance degradation in single-modality generation, state-of-the-art architectures employ progressive curriculum training combined with random modality masking.[5, 6] For example, the Klear framework is trained using a three-stage progressive strategy [5]:

- Stage-I: Pre-training: The model is trained on a large-scale, multi-scene corpus to acquire foundational generation capabilities.[5] This phase focuses on learning basic temporal synchronization, semantic alignment, and high-fidelity representation learning for both modalities individually.[5]
- Stage-II: Specialized Post-training: Guided by targeted evaluation metrics, the training data distribution is adaptively rebalanced to strengthen underperforming capabilities.[5] This stage focuses on improving weaker model functions, such as complex motion physics or fine-grained acoustic synthesis, without degrading general performance.[5]
- Stage-III: Quality-Refined Post-training: Finally, the model is fine-tuned on a high-quality, manually curated dataset to enhance perceptual realism, detail preservation, and instruction-following robustness.[5]

To ensure the model can perform joint generation (text-to-audio-video) and single-modality generation (text-to-video or text-to-audio) with equal proficiency, random modality masking is applied throughout all three training stages.[5, 6] By dynamically masking either the visual or auditory channel during training, the model is forced to develop independent representation capabilities while maintaining the ability to fuse both modalities when they are simultaneously active.[5, 6]

Automated Data Curation at Scale

The performance of unified models depends heavily on the volume and quality of their training data.[5, 6] To address the scarcity of high-quality, dense-captioned audiovisual data, modern systems utilize automated pipelines to curate large datasets.[6, 26]

These automated pipelines process millions of raw video clips through several stages:

By applying this automated data strategy, Klear scales its training corpus to 81 million samples with dense captions [26], while MM-Sonate leverages a pre-training corpus of 100 million strictly aligned audio-video-caption triplets.[25]

For community reproducibility, researchers have released targeted paired datasets, including 13 hours of video-game clips and 64 hours of concert performances, each segmented into consistent 34-second samples.[24, 27] These datasets serve as standardized testbeds to evaluate joint architectures, demonstrating their ability to synthesize semantically coherent audio-video pairs and quantitatively assess alignment during rapid on-screen actions and complex musical cues.[24, 27]

High-Value Production Workflows and Temporal Synchronization

Integrating unified models into professional production workflows requires precise control over character consistency, camera movement, and sound design.[28, 29, 30] The evolution of the Seedance framework highlights these advancements, showing a clear shift from the single-shot generation of Seedance 1.5 Pro to the multi-shot, reference-driven generation of Seedance 2.0.[28, 29, 30, 31]

Multi-Shot Continuity and Multimodal Referencing

Seedance 1.5 Pro, built on a dual-branch DiT architecture, generated synchronized audio and video in a single pass but struggled with multi-shot consistency, often treating camera instructions as loose suggestions and experiencing narrative drift across cuts.[28, 30] Seedance 2.0 addresses these limitations by utilizing a unified diffusion transformer that models relational dependencies across frames, enabling stable multi-shot generation with sequence-level planning.[30] Rather than treating shots independently, the model tracks characters, environments, and lighting across transitions, preventing random resets or visual drifting.[30] Additionally, Seedance 2.0 supports quad-modal inputs, allowing creators to upload up to 9 image references, 3 video references, and 3 audio references simultaneously.[26, 30] In this workflow, reference inputs serve as structural anchors: image references define consistent character identities and art direction; video references guide complex motion paths and camera choreography; and audio references shape dialog timing, sound effects, and spatial acoustic design.[29, 30]

Feature Dimension	Seedance 1.5 Pro	Seedance 2.0	Primary Source
Official Positioning	Native audio-visual joint generation model [29]	Next-generation model with reference-driven creation and enhanced control [29]	[29]

Max Resolution	1080p (T2V/I2V support) [28, 31]	1080p Native (Up to 2K supported) [30, 32]	[28, 31, 32]
Max Generation Duration	12 seconds [31]	15 seconds [31]	[31]
Reference Capacity	Single reference image support [29]	Quad-modal (Up to 9 images, 3 videos, 3 audio segments per request) [26, 30]	[26, 29, 30]
Audio Quality & Sync	Multilingual speech and spatial sound effects [29]	Immersive audio-video integration, improved lip-sync, richer sound design [29]	[29]
Camera Movement & Composition	Cinematic camera movement control [28]	Precise motion control, flexible camera choreography, Dutch angle, Dolly Zoom [32]	[28, 32]
Typical Use Cases	Short-form entertainment, simple localized ads [28, 29]	Professional commercials, multi-shot series, brand campaigns [29, 32]	[28, 29, 32]

Decoupling Speaker Identity via Timbre Injection

For dialogue-heavy scenes, maintaining speaker consistency across shots is critical.[15, 33] In MM-Sonate, this is achieved through a specialized parameter-efficient timbre injection mechanism that decouples speaker identity from the linguistic content of the phoneme stream.[15, 16, 17] During generation, an audio reference containing the target speaker's voice is processed to extract a dense timbre embedding.[16, 17]

This embedding is injected directly into the joint MM-DiT block.[16] By conditioning the cross-modal attention layers on this timbre vector, the model performs zero-shot voice cloning.[15, 16, 17] It synthesizes synchronized speech and lip motion matching the reference

speaker's voice, while ensuring the corresponding video frames maintain high visual identity preservation.[16, 17]

To further enhance audio realism, MM-Sonate replaces the standard zero-vector negative conditioning used in Classifier-Free Guidance with a noise-based negative conditioning strategy.[15, 17, 25] By using natural background noise as a negative prior during sampling, the model is guided away from synthetic artifacts and acoustic distortion, yielding clearer speech and more natural ambient room tone.[15, 17, 25]

Real-time Streaming Bottlenecks

While joint models achieve remarkable quality, their bidirectional attention dependencies introduce high latency, hindering interactive and real-time application deployment.[19] To address this temporal bottleneck, several architectural solutions have been introduced:

- OmniForcing: This framework optimizes joint audio-visual diffusion models for real-time applications by restructuring attention dependencies, significantly reducing the latency associated with bidirectional cross-attention passes.[19]
 - Hallo-Live: Designed for text-driven joint avatar generation, this streaming framework synthesizes portrait video and synchronized speech by combining asynchronous dual-stream diffusion with human-centric preference-guided distillation, ensuring high fidelity under aggressive acceleration.[19]
 - Asynchronous Noise and Motion Schedules: Backbones such as A2V-DiT utilize asynchronous add-noise and asynchronous motion-guided generation in the latent space, maintaining consistency in generated video clips while allowing streaming motion generation in real-time exceeding 30 frames per second.[18, 34]
-

Continuous Latent Space Analytical Processing and Multimodal Deduction

As unified models mature, their role is expanding beyond passive content generation to include active multimodal comprehension and continuous latent-space deduction.[35, 36]

The Bottleneck of Explicit Textual Representation

Traditional Multimodal Large Language Models (MLLMs) perform step-by-step cognitive tasks by converting sensory inputs (such as video frames and audio clips) into discrete text tokens.[35, 36] However, this explicit text-based deduction process introduces significant limitations [35, 36, 37]:

- Information Bottleneck: Highly detailed, high-dimensional sensory signals—including subtle acoustic pitch shifts, complex physical textures, and nuanced visual motion—are compressed into discrete text descriptions, losing critical temporal and spatial details.[35, 36, 37]

- Semantic Grounding Loss: Once mapped to text, the model's analytical steps shift toward language priors, reducing its attention to the original sensory evidence.[35, 36]
- Temporal Disconnection: Text tokens lack intrinsic temporal dimensions, making it difficult to maintain precise cross-modal coordination during multi-step inference.[35, 36]

To resolve these limitations, the LatentOmni framework performs step-by-step cognitive processing directly within a unified, continuous latent space.[35, 36, 38] Instead of relying on explicit text-based rationalizations, LatentOmni interleaves textual analysis with continuous audio-visual latent states.[35, 36, 38] This design allows the model to preserve dense, uncompressed sensory representations and remain grounded in the original physical inputs throughout the analytical trajectory.[35, 36, 38]

Omni-Sync Position Embedding (OSPE)

To ensure that the audio and video latent streams remain aligned during multi-step inference, LatentOmni introduces Omni-Sync Position Embedding (OSPE).[36, 38] OSPE extends standard time-aligned multimodal Rotary Position Embeddings (RoPE) to synchronized latent features.[35, 36]

Let h

i

v

and h

j

a

represent the latent visual and auditory feature representations at temporal indices i and j , respectively.[36] OSPE applies a synchronized rotation operator R

Θ

to these representations, ensuring that elements occurring at the same physical time share matching temporal coordinates within the attention space [35, 36]:

h

$\sim$

i

v

$=R$

Θ

$(i)h$

i

v

h

$\sim$

j

a

$=R$

Θ

$(j)h$

j

a

By synchronizing the position embeddings across different modalities, OSPE prevents temporal drift during long-context inference, enabling the model to precisely align auditory events with their visual counterparts.[35, 36, 38]

To train this architecture, the LatentOmni pipeline incorporates feature-level supervision.[35, 36, 38] It minimizes a joint reconstruction loss that encourages intermediate latent cognitive states to retain task-relevant visual and auditory features:

L

total

$=L$

text

$+\lambda$

1

L

recon_v

$+\lambda$

2

L

recon_a

This multi-task loss is optimized using a specialized dataset of interleaved audio-visual trajectories, such as LatentOmni-Instruct-35K, which contains question-answer pairs filtered for analytical difficulty, logical soundness, and modality dependency.[36, 38] This allows the model to perform continuous multimodal deduction while remaining grounded in the raw sensory inputs.[35, 36, 38]

Multi-Branch Adversarial Diagnostics and Evaluation Metrics

The rapid development of joint generative architectures has introduced significant risks, particularly regarding the synthesis of highly realistic, deepfake audiovisual content and potential semantic distortion.[39] This has necessitated the development of advanced evaluation benchmarks and multi-branch adversarial diagnostics.[19, 39]

Minute-Scale Evaluation Benchmarks

Traditional evaluation protocols for generative models are largely confined to short clips, failing to capture how identity consistency, narrative coherence, and audio-visual alignment degrade over extended temporal horizons.[19] To address this evaluation gap, the LongAV-Compass benchmark establishes a diagnostic framework specifically designed for minute-scale audio-visual generation across text-to-audiovisual, image-to-audiovisual, and video-to-audiovisual tasks.[19]

LongAV-Compass evaluates 20 fine-grained dimensions categorizing generation quality into three primary tiers:

- **Within-Segment Quality:** Focuses on immediate visual fidelity, acoustic clarity, and localized lip-sync precision within individual scenes.[19]
- **Cross-Segment Consistency:** Assesses the continuity of character appearances, ambient sound design, and lighting styles across scene transitions.[19]
- **Global Narrative Coherence:** Monitors the long-term thematic alignment and structural progression of the generated audio-visual sequence over several minutes, ensuring that temporal progression does not result in semantic drift.[19]

Multi-Branch Adversarial Detection Architectures

To detect and localize partially manipulated AI-generated video forgeries across both visual and audio channels, modern security frameworks utilize multi-branch architectures.[39] These detection models integrate complementary diagnostic streams to identify fine-grained, cross-modal inconsistencies.[39]

The core architecture of these systems is divided into three specialized branches:

- **Large Multimodal Model (LMM) Semantic Branch:** Analyzes high-level semantic and language-driven cues to expose conceptual contradictions, such as a speaker's mouth movements indicating a different word than the synthesized audio track.[39]
- **Spatio-Temporal (ST) Visual Branch:** Monitors pixel-level artifacts, unnatural physical motion, and temporal flickering across video frames.[39]
- **Multi-Scale Partial-Spoof (PS) Audio Branch:** Analyzes the acoustic track over multiple temporal resolutions, outputting segment-level anomaly scores to capture localized spoofing traces.[39]

By fusing these modality-specific prediction streams, the system produces dense, time-aligned likelihood streams.[39] This enables security frameworks to perform fine-grained temporal localization, identifying not only whether a video has been manipulated, but also pinpointing the exact temporal boundaries and designating which modality has been forged.[39] This multi-branch approach significantly outperforms traditional single-modality detection systems, illustrating that defending against generative deepfakes requires the same unified, multimodal understanding that drives advanced synesthetic generation.[39]

-
1. Diffusion Models for Joint Audio-Video Generation - arXiv, <https://arxiv.org/html/2603.16093v1>
 2. Seeing and Hearing: Open-domain Visual-Audio Generation with Diffusion Latent Aligners, https://openaccess.thecvf.com/content/CVPR2024/papers/Xing_Seeing_and_Hearing_Open-domain_Visual-Audio_Generation_with_Diffusion_Latent_Aligners_CVPR_2024_paper.pdf
 3. Seeing and Hearing: Open-domain Visual-Audio Generation with Diffusion Latent Aligners, <https://arxiv.org/html/2402.17723v1>
 4. AV-DiT: Efficient Audio-Visual Diffusion Transformer for Joint Audio and Video Generation - arXiv, <https://arxiv.org/html/2406.07686v1>
 5. Unified Multi-Task Audio-Video Joint Generation - arXiv, <https://arxiv.org/html/2601.04151v1>
 6. Klear: Unified Multi-Task Audio-Video Joint Generation - ResearchGate, https://www.researchgate.net/publication/399559825_Klear_Unified_Multi-Task_Audio-Video_Joint_Generation
 7. Synesthesia of Vehicles: Tactile Data Synthesis from Visual Inputs - arXiv, <https://arxiv.org/html/2602.01832v1>
 8. Robot Synesthesia: A Sound and Emotion Guided Robot Painter - arXiv, <https://arxiv.org/html/2302.04850v3>
 9. Stochastic resonance model of synaesthesia - Royal Society Publishing, <https://royalsocietypublishing.org/rstb/article/374/1787/20190029/30588/Stochastic-resonance-model-of>
 10. “Mickey Mousing” in the Brain: Motion-Sound Synesthesia and the Subcortical Substrate of Audio-Visual Integration - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7917298/>
 11. The Anatomy of Veo 3: DeepMind's Audiovisual Diffusion Model | by ..., <https://medium.com/@frinktyler1445/the-anatomy-of-veo-3-deepminds-audiovisual-diffusion-model-1721bec4b156>
 12. Align your Latents: High-Resolution Video Synthesis with Latent Diffusion Models, <https://research.nvidia.com/labs/toronto-ai/VideoLDM/>
 13. AV-DiT: Efficient Audio-Visual Diffusion Transformer for Joint Audio and Video Generation - OpenReview, <https://openreview.net/pdf/38547b90584348f68737c2c7f1f86a7623dfdb35.pdf>
 14. Training-Free Multimodal Guidance for Video to Audio Generation - arXiv, <https://arxiv.org/html/2509.24550v1>
 15. Chunyu Qiang - CatalyzeX, <https://www.catalyzex.com/author/Chunyu%20Qiang>

16. MM-Sonate: Multimodal Flow-Matching Transformer - Emergent Mind,
<https://www.emergentmind.com/topics/mm-sonate>
17. MM-Sonate: Multimodal Controllable Audio-Video Generation with Zero-Shot Voice Cloning, <https://arxiv.org/html/2601.01568v2>
18. READ: Real-time and Efficient Asynchronous Diffusion for Audio-driven Talking Head Generation | Proceedings of the AAAI Conference on Artificial Intelligence,
<https://ojs.aaai.org/index.php/AAAI/article/view/37940>
19. Daily Papers - Hugging Face,
<https://huggingface.co/papers?q=audio-visual%20generation>
20. Diff-Foley: Synchronized Video-to-Audio Synthesis with Latent Diffusion Models - arXiv,
<https://arxiv.org/html/2306.17203>
21. Diff-Foley: Synchronized Video-to-Audio Synthesis with Latent Diffusion Models | OpenReview, <https://openreview.net/forum?id=q5FAZAl0oz¬elId=aV652dfbzS>
22. DIFF-FOLEY: Synchronized Video-to-Audio Synthesis with Latent Diffusion Models,
https://proceedings.neurips.cc/paper_files/paper/2023/file/98c50f47a37f63477c01558600dd225a-Paper-Conference.pdf
23. GMS-CAVP: Improving Audio-Video Correspondence with Multi-Scale Contrastive and Generative Pretraining - arXiv, <https://arxiv.org/html/2601.19606v1>
24. Paper page - Diffusion Models for Joint Audio-Video Generation - Hugging Face,
<https://huggingface.co/papers/2603.16093>
25. MM-Sonate: Multimodal Controllable Audio-Video Generation with Zero-Shot Voice Cloning, <https://arxiv.org/html/2601.01568v1>
26. Paper page - Klear: Unified Multi-Task Audio-Video Joint Generation - Hugging Face,
<https://huggingface.co/papers/2601.04151>
27. [2603.16093] Diffusion Models for Joint Audio-Video Generation - arXiv,
<https://arxiv.org/abs/2603.16093>
28. Seedance 1.5 Pro by ByteDance | AI Model on Morpic,
<https://morphic.com/resources/models/seedance-1-5>
29. Seedance 2.0/1.5 Pro - AI Video Generator on xmk seedance2,
<https://www.xmk.com/seedance>
30. Seedance 2.0 vs 1.5 Pro: a direct comparison - Runware,
<https://runware.ai/blog/seedance-2-0-has-landed-on-runware>
31. Seedance 1.5 Pro vs Seedance 2.0 | What Changed, Upgrade... - MaxVideoAI,
<https://maxvideoai.com/ai-video-engines/seedance-1-5-pro-vs-seedance-2-0>
32. Seedance 2.0/1.5 Pro - AI Video Generator on jxp seedance2,
<https://www.jxp.com/seedance>
33. MM-Sonate: Multimodal Controllable Audio-Video Generation with Zero-Shot Voice Cloning,
https://www.researchgate.net/publication/399478724_MM-Sonate_Multimodal_Controllable_Audio-Video_Generation_with_Zero-Shot_Voice_Cloning
34. Daily Papers - Hugging Face,
<https://huggingface.co/papers?q=language-motion%20latent%20space>
35. Hao Liang - CatalyzeX, <https://www.catalyzex.com/author/Hao%20Liang>

36. LatentOmni: Rethinking Omni-Modal Understanding via Unified Audio-Visual Latent Reasoning - arXiv, <https://arxiv.org/html/2605.22012v1>
37. Yue Ding's research works - ResearchGate, <https://www.researchgate.net/scientific-contributions/Yue-Ding-2303257377>
38. yfanDai/LatentOmni - GitHub, <https://github.com/yfanDai/LatentOmni>
39. Towards Multi-Modal Forgery Representation Learning for AI-Generated Video Detection and Localization - arXiv, <https://arxiv.org/html/2605.07232v1>